

CSE185

Introduction to Computer Vision

Lab 03: Spatial Filters

Instructor: Daniel Leung
TA: Mohammadkazem Ebrahimpour
Xueqing Deng

Overviews

- Median filter:



Noisy input image



Median filter output

Overviews

- Sobel filter:



Input image




Sobel filter output

Overviews

- Gaussian filter:




Input image



Gaussian filter output

Spatial Filter

- A **spatial filter** is an image operation where each pixel value $I(u, v)$ is changed by a function of the intensities of pixels in a neighborhood of (u, v) .




- Mean filter:

$$I'(u, v) = \frac{1}{9} \sum_{i=-1}^1 \sum_{j=-1}^1 I(u + i, v + j)$$

Linear Filter

- H is the filter matrix (kernel):

$$I'(u, v) = \sum_{i=-1}^1 \sum_{j=-1}^1 I(u+i, v+j) \cdot H(i, j)$$



Linear Filter

- Assume H is a 3×3 mean filter: $H = \frac{1}{9} \begin{pmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{pmatrix}$



```
I1 = im2double(imread('lena.jpg'));  
I2 = zeros(size(I1));
```

```
for u = 1 : size(I1, 2)  
    for v = 1 : size(I1, 1)
```

```
        value = ???;  
        I2(v, u) = value;
```

```
    end
```

```
end
```



Linear Filter

$$I'(u, v) = \sum_{i=-1}^1 \sum_{j=-1}^1 I(u+i, v+j) \cdot H(i, j)$$

```
I1 = im2double(imread('lena.jpg'));  
I2 = zeros(size(I1));
```


```
for u = 1 : size(I1, 2)  
    for v = 1 : size(I1, 1)
```

```
        value = 0;  
        for i = -1:1  
            for j = -1:1  
                value = value + ???  
            end  
        end  
        I2(v, u) = value;
```

```
    end
```

```
end
```

Select neighborhood




Linear Filter

$$I'(u, v) = \sum_{i=-1}^1 \sum_{j=-1}^1 I(u+i, v+j) \cdot H(i, j)$$

```
I1 = im2double(imread('lena.jpg'));  
I2 = zeros(size(I1));
```

```
for u = 1 : size(I1, 2)  
    for v = 1 : size(I1, 1)  
  
        value = 0;  
        for i = -1:1  
            for j = -1:1  
                value = value + ???  
            end  
        end  
        I2(v, u) = value;  
    end  
end
```

Select neighborhood



Linear Filter

$$I'(u, v) = \sum_{i=-1}^1 \sum_{j=-1}^1 I(u+i, v+j) \cdot H(i, j)$$


```
I1 = im2double(imread('lena.jpg'));  
I2 = zeros(size(I1));
```

```
for u = 1 : size(I1, 2)  
    for v = 1 : size(I1, 1)
```


```
        value = 0;  
        for i = -1:1  
            for j = -1:1  
                value = value + ???  
            end  
        end  
        I2(v, u) = value;
```

```
    end
```

```
end
```



×



Linear Filter

$$I'(u, v) = \sum_{i=-1}^1 \sum_{j=-1}^1 I(u+i, v+j) \cdot H(i, j)$$

```
I1 = im2double(imread('lena.jpg'));  
I2 = zeros(size(I1));
```


```
for u = 1 : size(I1, 2)  
    for v = 1 : size(I1, 1)
```

```
        value = 0;  
        for i = -1:1  
            for j = -1:1  
                value = value + ???  
            end  
        end  
        I2(v, u) = value;
```

```
    end
```

```
end
```

Select neighborhood



Linear Filter

- Be careful around boundaries
 - ignore the pixels on boundaries
 - padding/extend boundaries [Optional]

```
I1 = im2double(imread('lena.jpg'));  
I2 = zeros(size(I1));
```

```
for u = 1 : size(I1, 2)  
    for v = 1 : size(I1, 1)
```

```
        value = 0;  
        for i = -1:1  
            for j = -1:1  
                value = value + ???
```


```
            end
```

```
        end
```

```
        I2(v, u) = value;
```

```
    end
```

```
end
```



Be careful about the index

Linear Filter

- Note: For Loop in MATLAB is very slow!
- Tips: use `tic` and `toc` to measure the elapsed time

```
%% Gaussian filter
hsize = 11;
sigma = 4;
tic
I = gaussian_filter_slow(img, hsize, sigma);
toc
tic
I = gaussian_filter_fast(img, hsize, sigma);
toc
```

```
Elapsed time is 3.522677 seconds.
Elapsed time is 1.608707 seconds.
```

```
fx >>
```

Linear Filter

- Extract patch/sub-matrix, and do matrix/vector operation.
- Again, be careful about the index around image boundaries

```
I1 = im2double(imread('lena.jpg'));  
I2 = zeros(size(I1));
```

```
for u = 1 : size(I1, 2)  
    for v = 1 : size(I1, 1)
```


```
        x1 = ???; x2 = ???;  
        y1 = ???; y2 = ???;
```

```
        patch = I1(y1:y2, x1:x2);  
        % convert matrix to vector  
        % matrix/vector operations  
        value = ???;  
        I2(v, u) = value;
```

```
    end
```

```
end
```

Check index range



Linear Filter

- Skip boundary pixels
 - How many pixels to shift?

```
I1 = im2double(imread('lena.jpg'));
I2 = zeros(size(I1));

for u = 1+shift_u : size(I1, 2)-shift_u
    for v = 1+ shift_v : size(I1, 1)- shift_v

        x1 = ???; x2 = ???;
        y1 = ???; y2 = ???;
        patch = I1(y1:y2, x1:x2);
        % convert matrix to vector
        % matrix/vector operations
        value = ???;
        I2(v, u) = value;

    end
end
```


Shift depends on the filter size

Median Filter

- Extract patch with the same size as the filter
- Calculate the median value of the patch (Use `median()`)
- Fill in the median value to the output pixel



Noisy input image



Median filter output

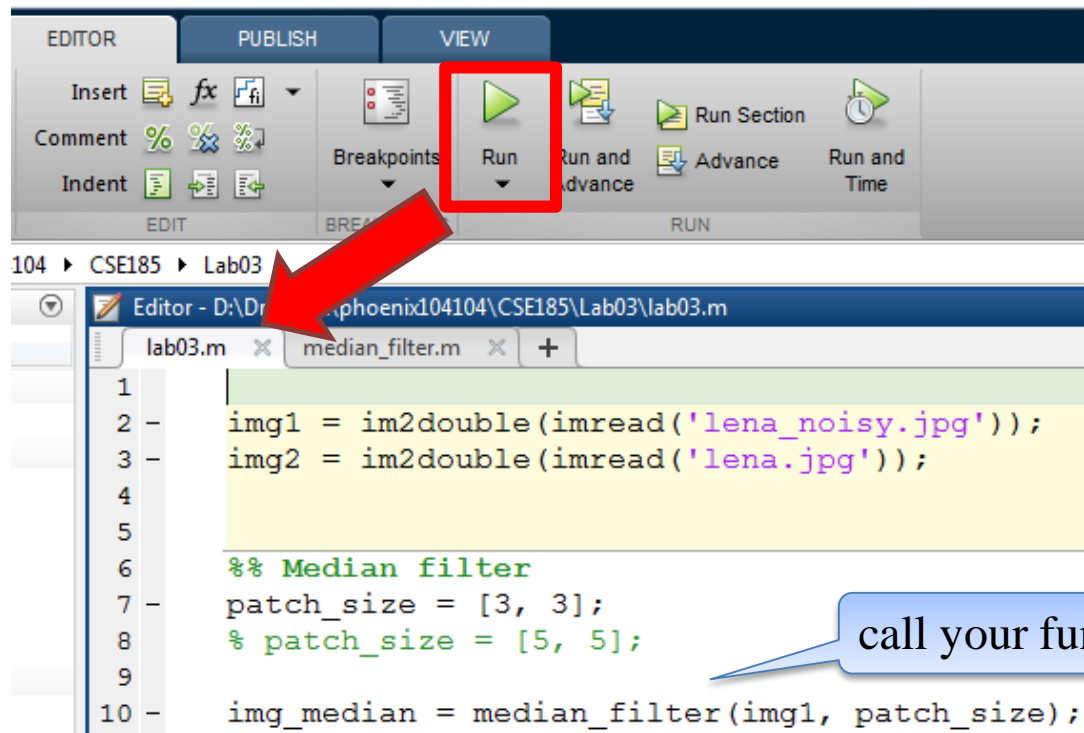
MATLAB Function

- Function template:

```
function output = median_filter(input, patch_size)
    output = input;
end
```

Assign output value, it will be automatically returned at the end of the function

- lab03.m is your main function, run this script instead of the function



call your function

Median Filter

- median_filter.m

```
function output = median_filter(img, patch_size)
    % YOUR CODE HERE
end
```

- In lab03.m:

```
img = im2double(imread('lena_noisy.jpg'));

%% Median filter
patch_size = [3, 3];
% patch_size = [5, 5];

img_median = median_filter(img, patch_size);
imwrite(img_median, 'median.jpg');
```

- Compare your result with built-in function:

```
I = medfilt2(img, patch_size);
```

Median Filter

- MATLAB function `median()`:


```
>> A = [1, 2, 3;  
        4, 5, 6;  
        7, 8, 9];  
>> median(A)  
ans =  
     4     5     6
```

- But we need a single value
 - convert matrix to vector first
 - or apply function twice
 - Question: are the results the same?

Sobel Filter

- Sobel filter is a simple edge detector

- $H = \begin{pmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{pmatrix}$ or $H = \begin{pmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{pmatrix}$



Input image

Sobel filter output

Sobel Filter

- `sobel_filter.m`

```
function output = sobel_filter(img, kernel)
    % YOUR CODE HERE
end
```

- In `lab03.m`:

```
img = im2double(imread('lena.jpg'));

%% Sobel filter
H = [1, 2, 1; 0, 0, 0; -1, -2, -1]; % horizontal edge
%H = [1, 0, -1; 2, 0, -2; 1, 0, -1]; % vertical edge

img_sobel = sobel_filter(img, H);
figure, imshow(img_sobel);
imwrite(img_sobel, 'sobel_h.jpg');
```

- Compare your result with `I = imfilter(img, H);`

Gaussian Filter

- Gaussian filter is a low-pass/smoothing filter

$$g(\Delta x, \Delta y) = \frac{1}{Z} \exp \left(-\frac{\Delta x^2 + \Delta y^2}{2\sigma^2} \right)$$



Input image

Gaussian filter output

Gaussian Filter

- Gaussian filter is a low-pass/smoothing filter

$$g(\Delta x, \Delta y) = \frac{1}{Z} \exp \left(-\frac{\Delta x^2 + \Delta y^2}{2\sigma^2} \right)$$

- $\Delta x, \Delta y$ are offsets from the kernel center, and Z is the normalized term to make the sum of all weights equal to 1
- In MATLAB:

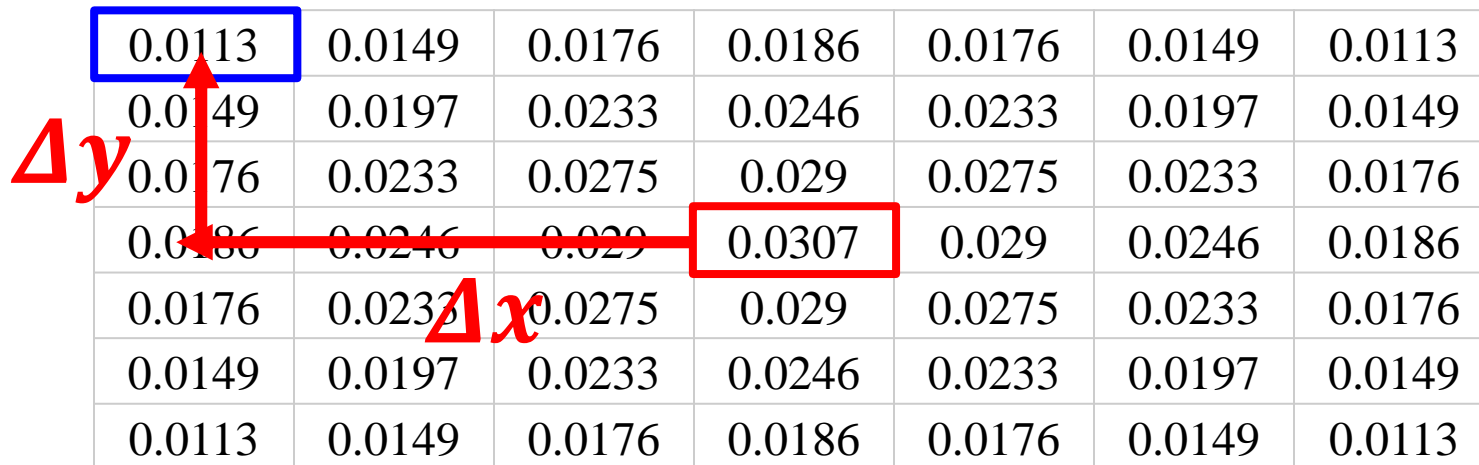
0.0113	0.0149	0.0176	0.0186	0.0176	0.0149	0.0113
0.0149	0.0197	0.0233	0.0246	0.0233	0.0197	0.0149
0.0176	0.0233	0.0275	0.029	0.0275	0.0233	0.0176
0.0186	0.0246	0.029	0.0307	0.029	0.0246	0.0186
0.0176	0.0233	0.0275	0.029	0.0275	0.0233	0.0176
0.0149	0.0197	0.0233	0.0246	0.0233	0.0197	0.0149
0.0113	0.0149	0.0176	0.0186	0.0176	0.0149	0.0113

Gaussian Filter

- Gaussian filter is a low-pass/smoothing filter

$$g(\Delta x, \Delta y) = \frac{1}{Z} \exp \left(-\frac{\Delta x^2 + \Delta y^2}{2\sigma^2} \right)$$

- $\Delta x, \Delta y$ are offsets from the kernel center, and Z is the normalized term to make the sum of all weights equal to 1
- In MATLAB:



A 7x7 Gaussian kernel matrix is shown. The center element (3,3) is 0.0307 and is highlighted with a red box. A red cross is drawn through the matrix, with a vertical line labeled Δy and a horizontal line labeled Δx indicating the offsets from the center. The top-left element (1,1) is 0.0113 and is highlighted with a blue box.

0.0113	0.0149	0.0176	0.0186	0.0176	0.0149	0.0113
0.0149	0.0197	0.0233	0.0246	0.0233	0.0197	0.0149
0.0176	0.0233	0.0275	0.029	0.0275	0.0233	0.0176
0.0186	0.0246	0.029	0.0307	0.029	0.0246	0.0186
0.0176	0.0233	0.0275	0.029	0.0275	0.0233	0.0176
0.0149	0.0197	0.0233	0.0246	0.0233	0.0197	0.0149
0.0113	0.0149	0.0176	0.0186	0.0176	0.0149	0.0113

Gaussian Filter


- You don't need to implement Gaussian kernel by yourself (this is bonus).

- Use `fspecial`:

```
H = fspecial('gaussian', hsize, sigma);
```

`hsize` is the size of the kernel, `sigma` = σ

- Visualization of a Gaussian kernel:



Gaussian Filter

- `gaussian_filter.m`

```
function output = gaussian_filter(img, hsize, sigma)
    H = fspecial('gaussian', hsize, sigma);
    % YOUR CODE HERE
end
```

- In `lab03.m`:

```
%% Gaussian filter
hsize = 5; sigma = 2;
% hsize = 9; sigma = 4;

img_gaussian = gaussian_filter(img, hsize, sigma);
figure, imshow(img_gaussian);
imwrite(img_gaussian, 'gaussian_5.jpg');
```

- Compare your result with `I = imfilter(img, H);`

TODO

1. Implement `median_filter.m`, use patch size = 3 and save the image as `median_3.jpg` (3pt)
2. Use patch size = 5, and save the image as `median_5.jpg` (3pt)
3. Implement `sobel_filter.m`, use horizontal filter and save the image as `sobel_h.jpg` (3pt)
4. Use vertical filter and save the image as `sobel_v.jpg` (3pt)
5. Implement `gaussian_filter.m`, use `hsize` = 5, `sigma` = 2, and save the image as `gaussian_5.jpg` (3pt)
6. Use `hsize` = 9, `sigma` = 4, and save the image as `gaussian_9.jpg` (3pt)
7. Upload your output images and `lab03.m` (2pt), `sobel_filter.m`, `median_filter.m`, and `gaussian_filter.m`

Bonus

- Implement boundary padding
 - zero padding (pixels outside the image are zero)
 - replicated/repeated padding (pixels outside the image are the same as pixels on the boundaries)
- You can try built-in function `padarray` or `wextend`, but you need to implement by yourself to get bonus points
- Implement Gaussian kernel
 - compute a matrix based on $g(\Delta x, \Delta y) = \exp\left(-\frac{\Delta x^2 + \Delta y^2}{2\sigma^2}\right)$
 - normalize the matrix to have sum equal to 1

Reference

- [Spatial Filtering](#)
- [Linear Algebra and MATLAB Tutorial](#)
- [Awesome Computer Vision](#)